

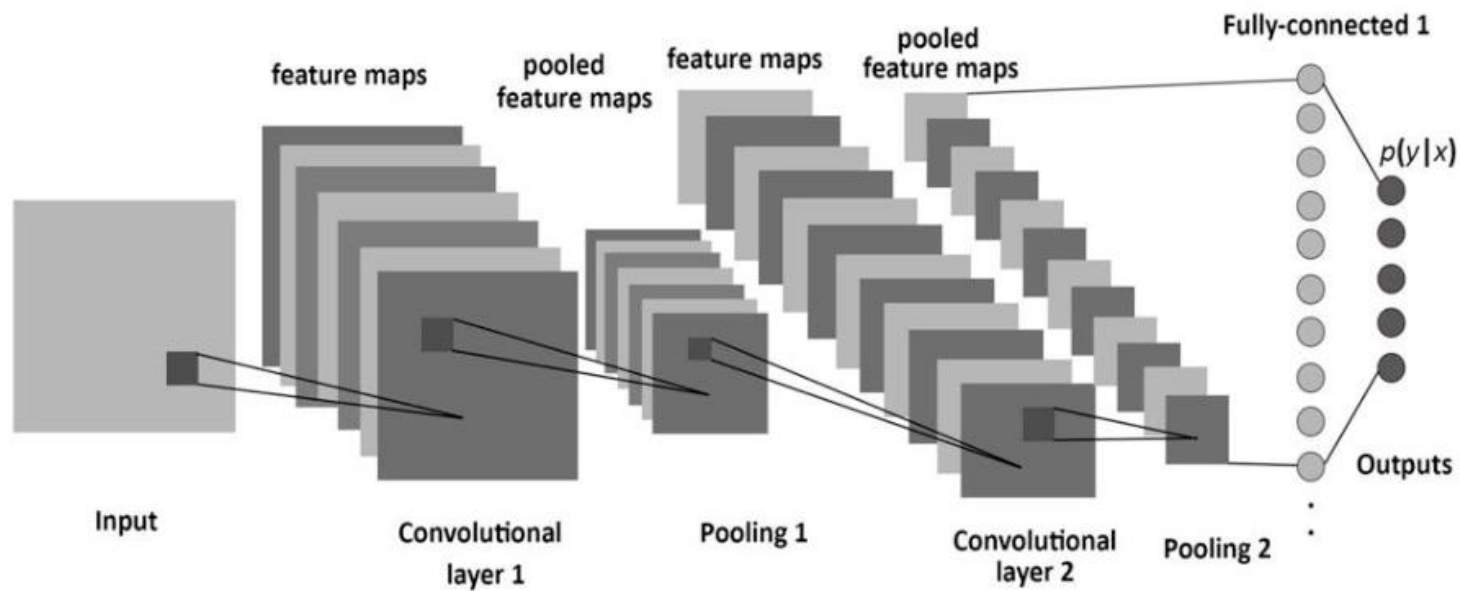
CNN

AILAB
Hanyang Univ.

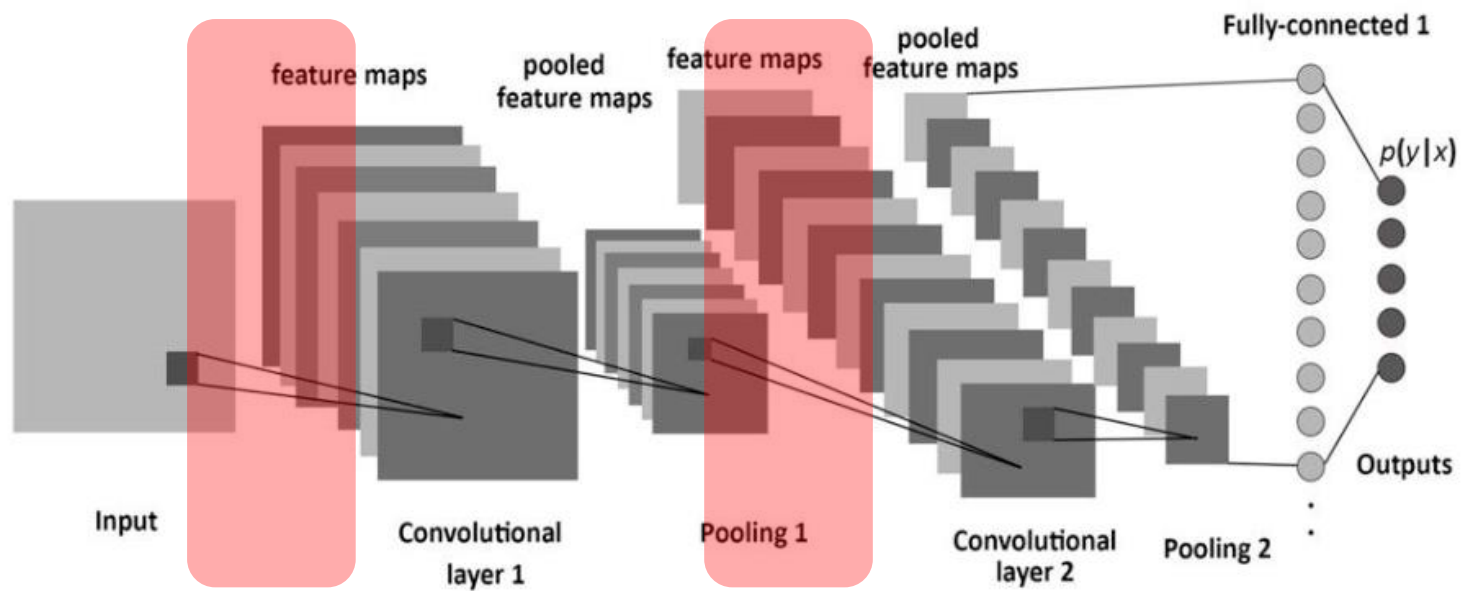
오늘 실습 내용

1. CNN 구현 MNIST Dataset
2. CNN 구현 CIFAR-10

CNN 구성요소 = (1) Convolution layer (2) Sub-sampling(pooling) layer (3) FC layer(Affine)



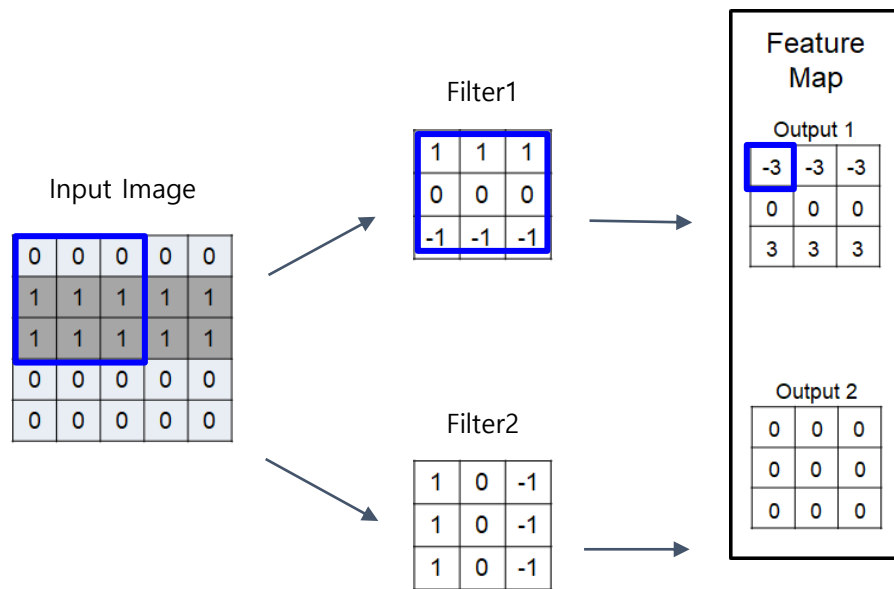
CNN 구성요소 = (1) Convolution layer (2) Sub-sampling(pooling) layer (3) FC layer(Affine)



Convolution

Convolution?

: Convolution(합성곱)은 하나의 함수와 또 다른 함수를 반전 이동한 값을 곱한 다음, 구간에 대해 적분하여 새로운 함수를 구하는 수학 연산자이다.



Convolution

torch.nn.Conv2d

<https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html>

```
torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0,  
dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None)
```

Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size (N, C_{in}, H, W) and output $(N, C_{\text{out}}, H_{\text{out}}, W_{\text{out}})$ can be precisely described as:


$$\text{out}(N_i, C_{\text{out}_j}) = \text{bias}(C_{\text{out}_j}) + \sum_{k=0}^{C_{\text{in}}-1} \text{weight}(C_{\text{out}_j}, k) \star \text{input}(N_i, k)$$

where \star is the valid 2D **cross-correlation** operator, N is a batch size, C denotes a number of channels, H is a height of input planes in pixels, and W is width in pixels.


Convolution

Channel


- 입력 데이터의 Channel 수와 필터의 Channel 수가 일치 해야 함
- 필터의 개수가 아웃풋의 Channel을 결정



Input image
(3, 3, 1)
(width, height, channel)



filter
(2, 2, 1)
(width, height, channel)




output
(2, 2, 1)
(width, height, channel)

Convolution

Channel


- 입력 데이터의 Channel 수와 필터의 Channel 수가 일치 해야 함
- 필터의 개수가 아웃풋의 Channel을 결정



Convolution

Channel

- 입력 데이터의 Channel 수와 필터의 Channel 수가 일치 해야 함
- 필터의 개수가 아웃풋의 Channel을 결정



Convolution

torch.nn.Conv2d

<https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html>

```
torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0,  
dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None)
```

Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size (N, C_{in}, H, W) and output $(N, C_{\text{out}}, H_{\text{out}}, W_{\text{out}})$ can be precisely described as:

$$\text{out}(N_i, C_{\text{out}_j}) = \text{bias}(C_{\text{out}_j}) + \sum_{k=0}^{C_{\text{in}}-1} \text{weight}(C_{\text{out}_j}, k) \star \text{input}(N_i, k)$$


where \star is the valid 2D **cross-correlation** operator, N is a batch size, C denotes a number of channels, H is a height of input planes in pixels, and W is width in pixels.

N : batch size
 C : # channel
 H : height
 W : width

Convolution

Channel

- 입력 데이터의 Channel 수와 필터의 Channel 수가 일치 해야 함
- 필터의 개수가 아웃풋의 Channel을 결정



Convolution

torch.nn.Conv2d

<https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html>

```
torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0,  
dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None)
```

Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size (N, C_{in}, H, W) and output $(N, C_{\text{out}}, H_{\text{out}}, W_{\text{out}})$ can be precisely described as:

$$\text{out}(N_i, C_{\text{out}_j}) = \text{bias}(C_{\text{out}_j}) + \sum_{k=0}^{C_{\text{in}}-1} \text{weight}(C_{\text{out}_j}, k) \star \text{input}(N_i, k)$$

where \star is the valid 2D **cross-correlation** operator, N is a batch size, C denotes a number of channels, H is a height of input planes in pixels, and W is width in pixels.

N : batch size
 C : # channel
 H : height
 W : width

Convolution

torch.nn.Conv2d

<https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html>

```
torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0,  
dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None)
```

Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size (N, C_{in}, H, W) and output $(N, C_{\text{out}}, H_{\text{out}}, W_{\text{out}})$ can be precisely described as:

$$\text{out}(N_i, C_{\text{out}_j}) = \text{bias}(C_{\text{out}_j}) + \sum_{k=0}^{C_{\text{in}}-1} \text{weight}(C_{\text{out}_j}, k) \star \text{input}(N_i, k)$$

where \star is the valid 2D cross-correlation operator, N is a batch size, C denotes a number of channels, H is a height of input planes in pixels, and W is width in pixels.

Convolution

torch.nn.Conv2d

<https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html>

```
torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0,  
dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None)
```

Stride

: 필터를 적용하는 위치의 간격

0	0	0	0	0
1	1	1	1	1
1	1	1	1	1
0	0	0	0	0
0	0	0	0	0

0	0	0	0	0
1	1	1	1	1
1	1	1	1	1
0	0	0	0	0
0	0	0	0	0

0	0	0	0	0
1	1	1	1	1
1	1	1	1	1
0	0	0	0	0
0	0	0	0	0

Stride = 1

0	0	0	0	0
1	1	1	1	1
1	1	1	1	1
0	0	0	0	0
0	0	0	0	0

0	0	0	0	0
1	1	1	1	1
1	1	1	1	1
0	0	0	0	0
0	0	0	0	0

Stride = 2

Convolution


torch.nn.Conv2d

<https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html>

```
torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0,
dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None)
```

Padding

: Convolution으로 인한 image 모서리 부분 정보 손실 방지를 위해 입력데이터 주변을 특정값으로 채우는 것 (일반적으로 Zero Padding 사용)



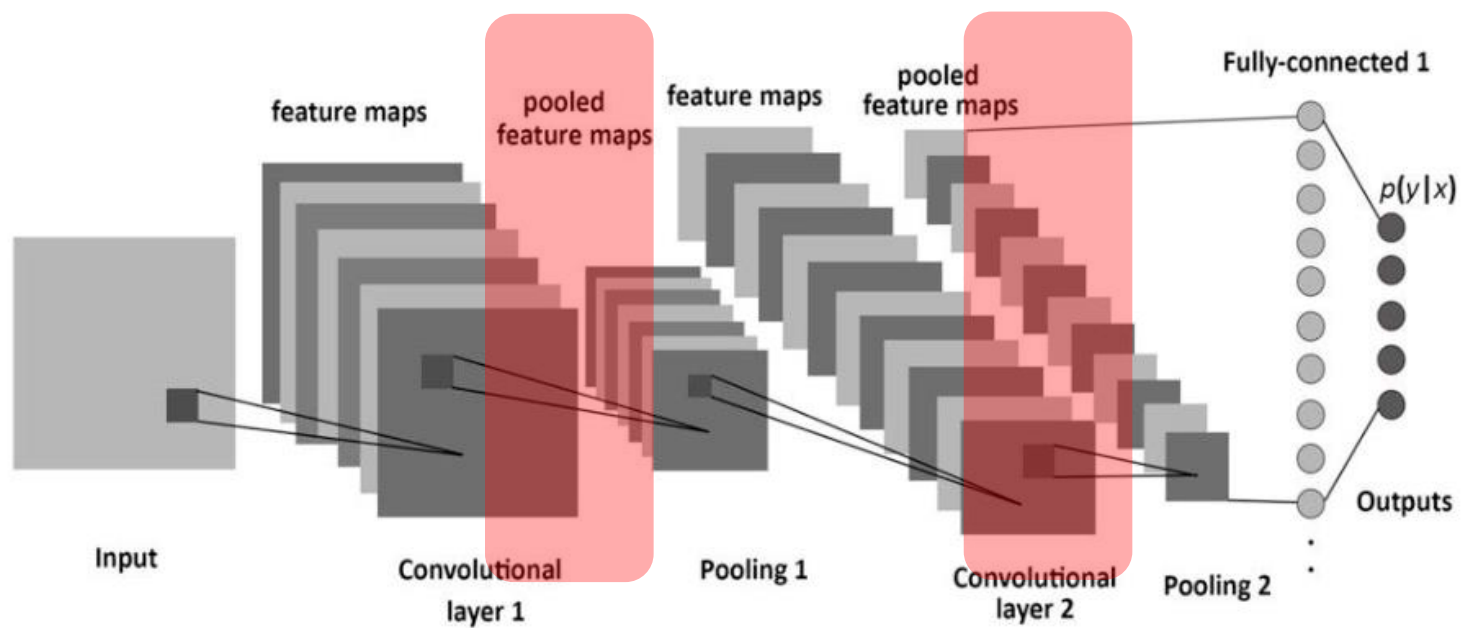
zero padding 사용X
7*7 → 5*5

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

zero padding 사용
7*7 → 7*7

`padding='valid'` is the same as no padding. `padding='same'` pads the input so the output has the shape as the input. However, this mode doesn't support any stride values other than 1.

CNN 구성요소 = (1) Convolution layer (2) Sub-sampling(pooling) layer (3) FC layer(Affine)





Pooling

Pooling?

: 가로 세로 방향의 공간을 줄이는 연산

- 출력의 해상도를 낮춰 변형이나 이동에 대한 민감도를 감소
- 이미지의 크기를 줄이기 때문에 학습할 노드의 수가 줄어들어 학습속도를 높이는 효과
- 하지만, 정보의 손실이 일어남
- CNN에서는 일반적으로 Max Pooling 사용



Pooling

torch.nn.MaxPool2d

<https://pytorch.org/docs/stable/generated/torch.nn.MaxPool2d.html#torch.nn.MaxPool2d>

```
torch.nn.MaxPool2d(kernel_size, stride=None, padding=0, dilation=1,  
return_indices=False, ceil_mode=False) [SOURCE]
```

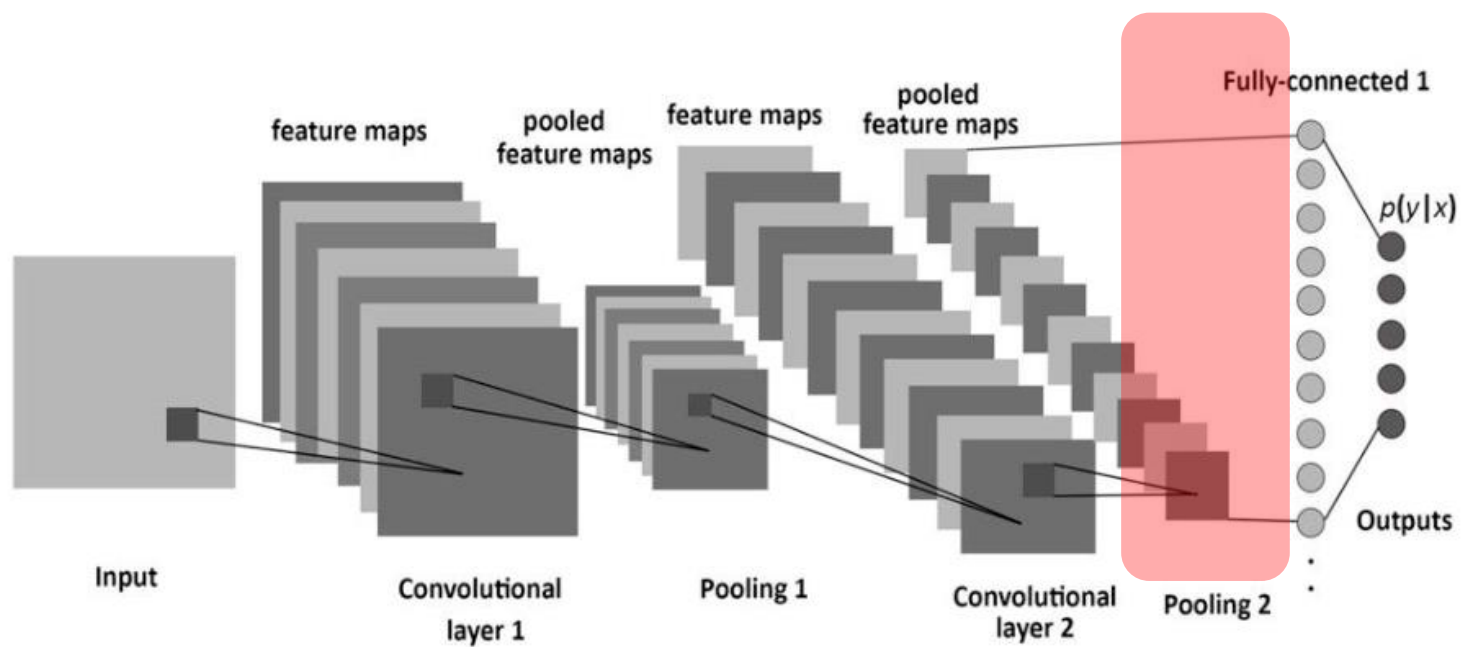
Applies a 2D max pooling over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size (N, C, H, W) , output (N, C, H_{out}, W_{out}) and `kernel_size` (kH, kW) can be precisely described as:

$$out(N_i, C_j, h, w) = \max_{m=0, \dots, kH-1} \max_{n=0, \dots, kW-1} \text{input}(N_i, C_j, \text{stride}[0] \times h + m, \text{stride}[1] \times w + n)$$

If `padding` is non-zero, then the input is implicitly padded with negative infinity on both sides for `padding` number

CNN 구성요소 = (1) Convolution layer (2) Sub-sampling(pooling) layer (3) FC layer(Affine)



Pooling

torch.flatten

<https://pytorch.org/docs/stable/generated/torch.flatten.html#torch.flatten>

nn.Linear를 사용하기 위해서 인풋을 반드시 1d로 만들어야 한다.

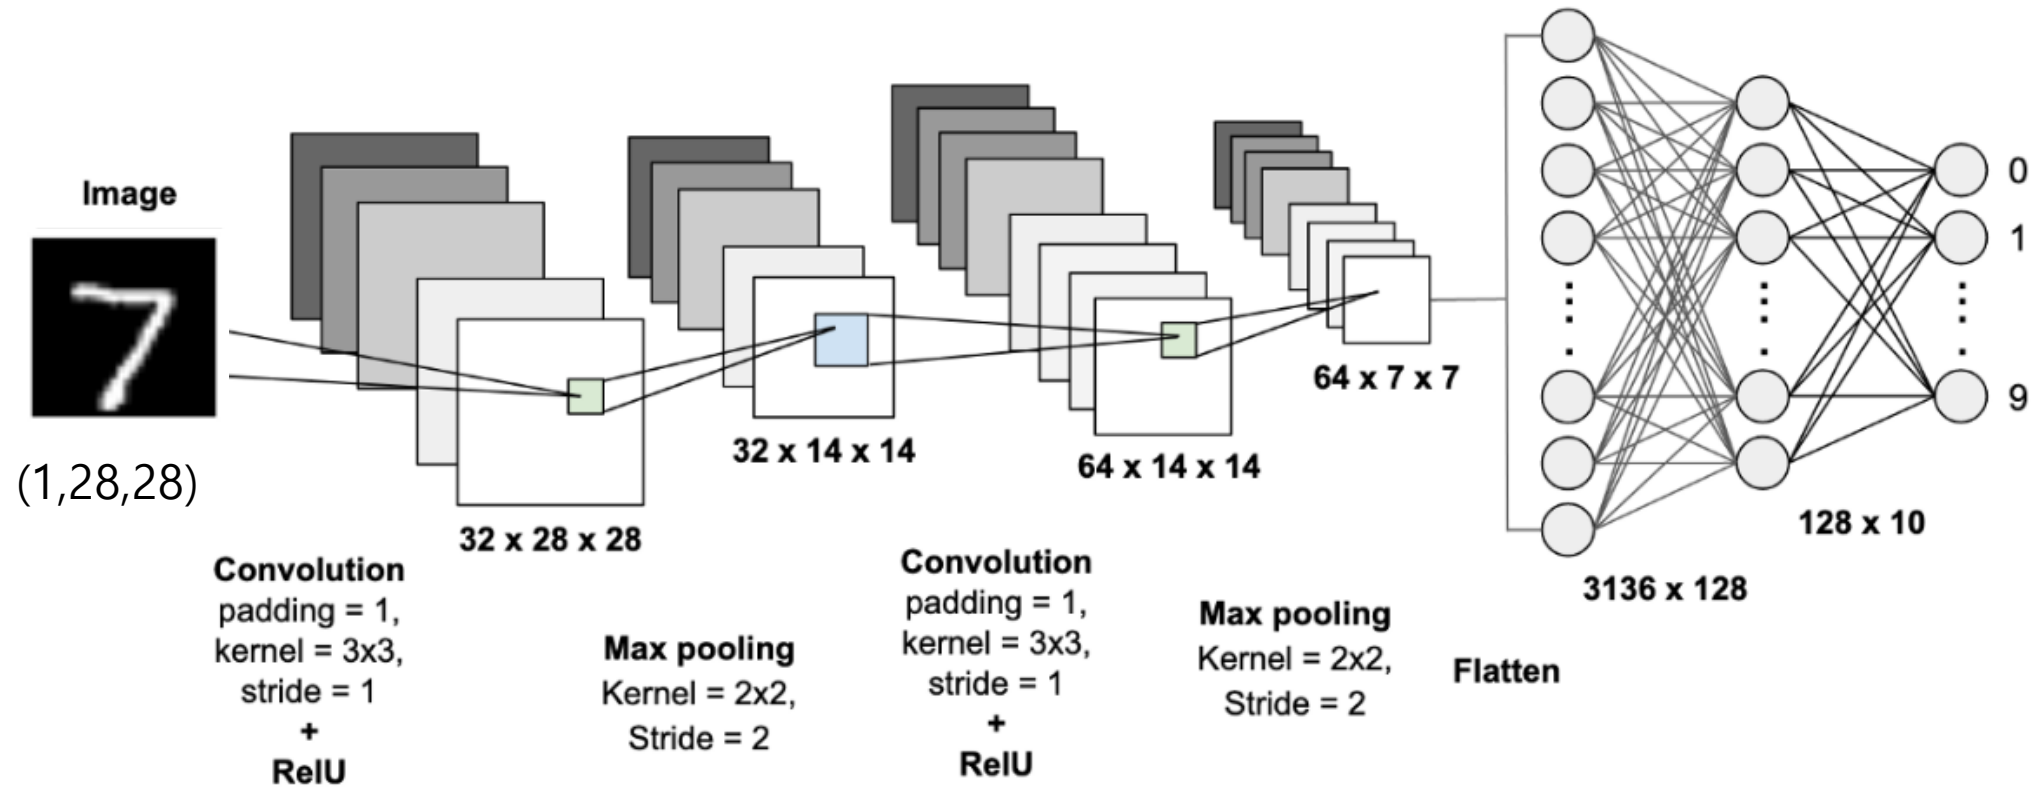
```
torch.flatten(input, start_dim=0, end_dim=- 1)
```

Flattens `input` by reshaping it into a one-dimensional tensor. If `start_dim` or `end_dim` are passed, only dimensions starting with `start_dim` and ending with `end_dim` are flattened. The order of elements in `input` is unchanged.

MNIST dataset

MNIST dataset

구조 출처: <https://towardsdatascience.com/mnist-handwritten-digits-classification-using-a-convolutional-neural-network-cnn-af5fafbc35e9>



MNIST dataset

```
import torch
import torch.nn as nn
import torch.optim as optim
import matplotlib.pyplot as plt
import numpy as np

torch.manual_seed(0)
torch.cuda.manual_seed(0)
torch.cuda.manual_seed_all(0)

if torch.cuda.is_available():
    device = torch.device('cuda')
else:
    device = torch.device('cpu')

import torchvision
import torchvision.transforms as transforms

train_dataset = torchvision.datasets.MNIST(root="MNIST_data/",
                                           train=True,
                                           transform=transforms.ToTensor(),
                                           download=True)
test_dataset = torchvision.datasets.MNIST(root="MNIST_data/",
                                           train=False,
                                           transform=transforms.ToTensor(),
                                           download=True)

batch_size = 128

train_dataloader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
test_dataloader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size)
```

```
x,y=train_dataset[0]
```

```
x.shape
```

```
torch.Size([1, 28, 28])
```

MNIST dataset

```
class CNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(1, 32, 3, 1, 1)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(32, 64, 3, 1, 1)
        self.fc1 = nn.Linear(64*7*7, 128)
        self.fc2 = nn.Linear(128, 10)
        self.activation = nn.ReLU()

    def forward(self, x):
        x = self.pool(self.activation(self.conv1(x)))
        x = self.pool(self.activation(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = self.activation(self.fc1(x))
        x = self.fc2(x)
        return x
```

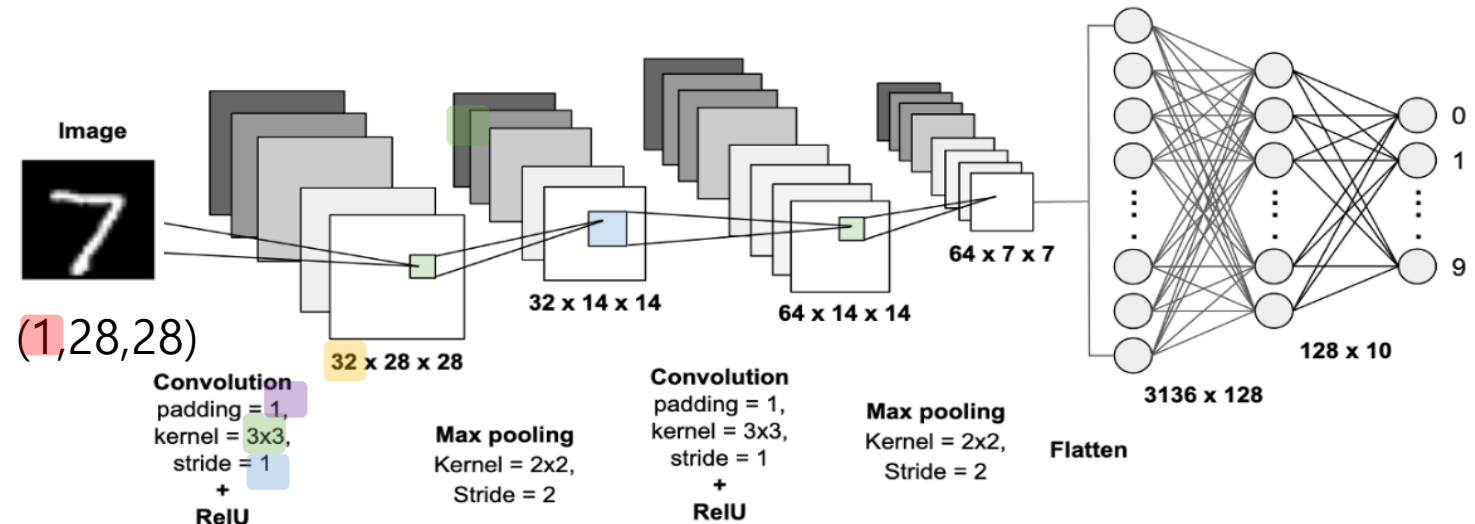
MNIST dataset

```

class CNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(1, 32, 3, 1, 1)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(32, 64, 3, 1, 1)
        self.fc1 = nn.Linear(64*7*7, 128)
        self.fc2 = nn.Linear(128, 10)
        self.activation = nn.ReLU()

    def forward(self, x):
        x = self.pool(self.activation(self.conv1(x)))
        x = self.pool(self.activation(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = self.activation(self.fc1(x))
        x = self.fc2(x)
        return x

```



MNIST dataset

```
model=CNN().to(device)

optimizer = optim.Adam(model.parameters(), lr=0.001) # set optimizer

criterion = nn.CrossEntropyLoss()

epochs = 30

model.train()
for epoch in range(epochs):
    model.train()
    avg_cost = 0
    total_batch_num = len(train_dataloader)

    for b_x, b_y in train_dataloader:
        logits = model(b_x.to(device)) # forward propagation
        loss = criterion(logits, b_y.to(device)) # get cost

    avg_cost += loss / total_batch_num
    optimizer.zero_grad()
    loss.backward() # backward propagation
    optimizer.step() # update parameters
    print('Epoch : {} / {}, cost : {}'.format(epoch+1, epochs, avg_cost))
```

MNIST dataset

```
correct = 0
total = 0

model.eval()
for b_x, b_y in test_dataloader:
    with torch.no_grad():
        logits = model(b_x.to(device))

        probs = nn.Softmax(dim=1)(logits)
        predicts = torch.argmax(logits, dim=1)

        total += len(b_y)
        correct += (predicts == b_y.to(device)).sum().item()

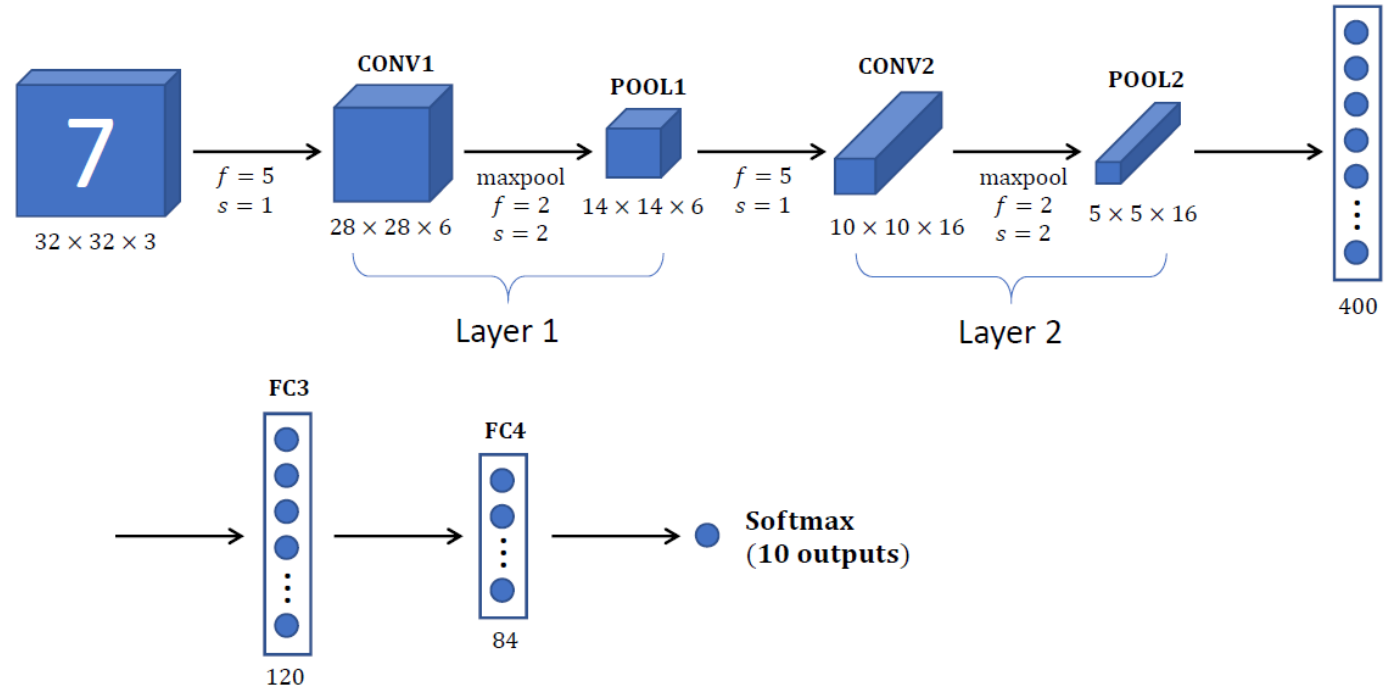
print(f'Accuracy of the network on test images: {100 * correct // total} %')
```

```
Accuracy of the network on test images: 99 %
```

CIFAR-10 dataset

 $(3, 32, 32)$

Neural network example (LeNet-5)



CIFAR-10 dataset

```
[1] import torch
import torch.nn as nn
import torch.optim as optim

[2] torch.manual_seed(0)
torch.cuda.manual_seed(0)
torch.cuda.manual_seed_all(0)

[3] if torch.cuda.is_available():
    device = torch.device('cuda')
else:
    device = torch.device('cpu')

[4] import torchvision
import torchvision.transforms as transforms

train_dataset = torchvision.datasets.CIFAR10(root="CIFAR10/",
                                              train=True,
                                              transform=transforms.ToTensor(),
                                              download=True)
test_dataset = torchvision.datasets.CIFAR10(root="CIFAR10/",
                                              train=False,
                                              transform=transforms.ToTensor(),
                                              download=True)

[5] batch_size = 128

train_dataloader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
test_dataloader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size)
```

```
x,y=train_dataset[0]
x.shape

torch.Size([3, 32, 32])
```

CIFAR-10 dataset

```
class CNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
        self.activation = nn.ReLU()

    def forward(self, x):
        x = self.pool(self.activation(self.conv1(x)))
        x = self.pool(self.activation(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = self.activation(self.fc1(x))
        x = self.activation(self.fc2(x))
        x = self.fc3(x)
        return x
```

CIFAR-10 dataset

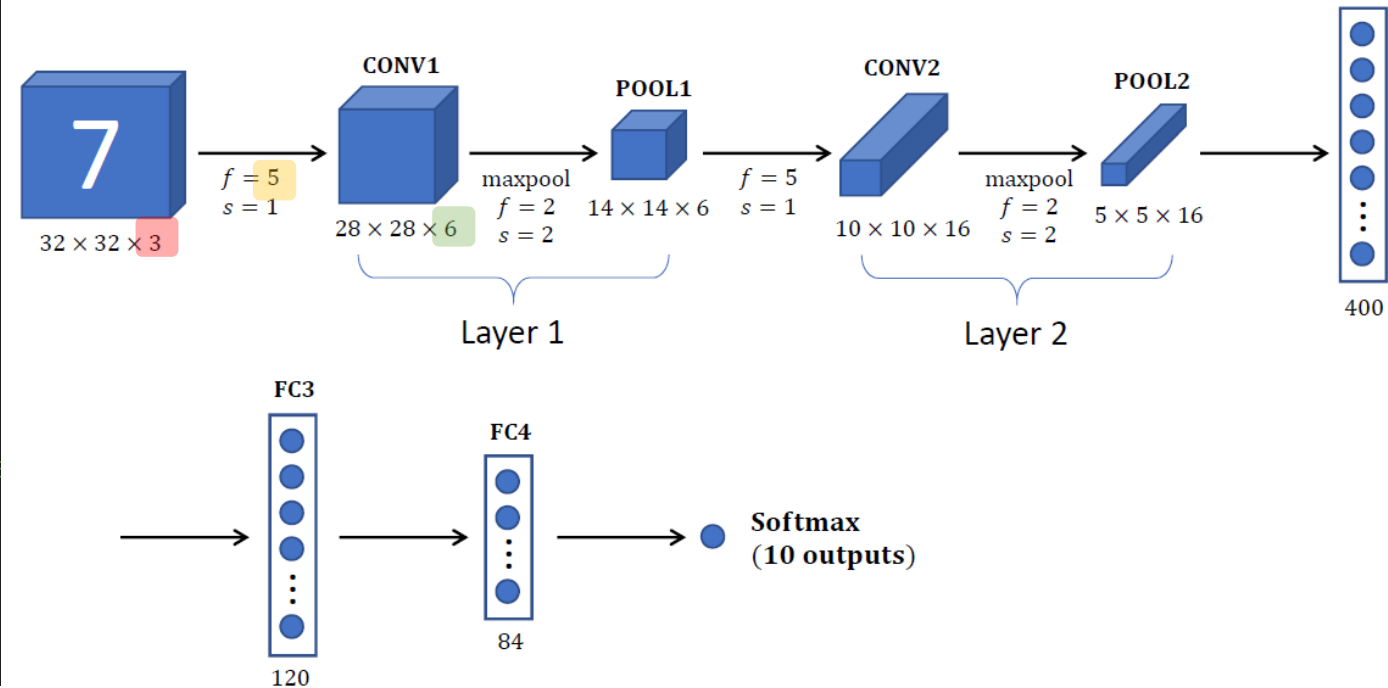
```

class CNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
        self.activation = nn.ReLU()

    def forward(self, x):
        x = self.pool(self.activation(self.conv1(x)))
        x = self.pool(self.activation(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = self.activation(self.fc1(x))
        x = self.activation(self.fc2(x))
        x = self.fc3(x)
        return x

```

Neural network example (LeNet-5)



CIFAR-10 dataset

```
model=CNN().to(device)

optimizer = optim.Adam(model.parameters(), lr=0.001) # set optimizer

criterion = nn.CrossEntropyLoss()

epochs = 30

model.train()
for epoch in range(epochs):
    model.train()
    avg_cost = 0
    total_batch_num = len(train_dataloader)

    for b_x, b_y in train_dataloader:
        logits = model(b_x.to(device)) # forward propagation
        loss = criterion(logits, b_y.to(device)) # get cost

        avg_cost += loss / total_batch_num
        optimizer.zero_grad()
        ▶ loss.backward() # backward propagation
        optimizer.step() # update parameters
    print('Epoch : {} / {}, cost : {}'.format(epoch+1, epochs, avg_cost))
```

CIFAR-10 dataset

```
correct = 0
total = 0

model.eval()
for b_x, b_y in test_dataloader:
    with torch.no_grad():
        logits = model(b_x.to(device))

        probs = nn.Softmax(dim=1)(logits)
        predicts = torch.argmax(logits, dim=1)

        total += len(b_y)
        correct += (predicts == b_y.to(device)).sum().item()

print(f'Accuracy of the network on test images: {100 * correct // total} %')
```

```
Accuracy of the network on test images: 63 %
```


오늘 실습 내용

1. MNIST에 대해 CNN구현
2. CIFAR-10에 대해 CNN구현